

Practical Machine Learning Project

Alan Zablocki

June 15, 2016

Setting up

We load the required libraries.

```
library(caret)
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
library(randomForest)
```

```
## randomForest 4.6-12
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##
```

```
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:ggplot2':
```

```
##
```

```
##     margin
```

```
library(rattle)
```

```
## Rattle: A free graphical interface for data mining with R.
```

```
## Version 4.1.0 Copyright (c) 2006-2015 Togaware Pty Ltd.
```

```
## Type 'rattle()' to shake, rattle, and roll your data.
```

```
library(rpart)
```

```
library(rpart.plot)
```

```
library(corrplot)
```

Load local data.

```
training = read.csv("pml-training.csv")
```

```
testing = read.csv("pml-testing.csv")
```

Check the dimensions.

```
dim(training)
```

```
## [1] 19622 160
```

```
dim(testing)
```

```
## [1] 20 160
```

Find out how many NA values using complete.cases()

```
sum(complete.cases(training))
```

```
## [1] 406
```

Cleaning the data

We will remove columns with NA values, columns with data that is not needed for training, e.g., names and time stamps etc.

We remove columns with NA values like this:

```
train_cleaned <- training[, colSums(is.na(training)) == 0]
```

There are seven columns at the beginning with information we do not need

```
sum(grepl("^X|user_name|timestamp|window", names(train_cleaned)))
```

```
## [1] 7
```

```
remove_train_col <- grepl("^X|user_name|timestamp|window", names(train_cleaned))
train_cleaned <- train_cleaned[, !remove_train_col]
ncol(train_cleaned)
```

```
## [1] 86
```

We are left with 86 columns but some have no values, and so we remove those too. Our “classe” variable will be removed but we column bind it back and make sure it is still called “class”.

```
trClean <- train_cleaned[, sapply(train_cleaned, is.numeric)]
trained <- cbind(trClean, train_cleaned[, "classe"])
names(trained)[53] <- "classe"
colnames(trained)
```

```
## [1] "roll_belt"          "pitch_belt"         "yaw_belt"
## [4] "total_accel_belt"   "gyros_belt_x"       "gyros_belt_y"
## [7] "gyros_belt_z"       "accel_belt_x"       "accel_belt_y"
## [10] "accel_belt_z"       "magnet_belt_x"      "magnet_belt_y"
## [13] "magnet_belt_z"      "roll_arm"           "pitch_arm"
```

```
## [16] "yaw_arm"          "total_accel_arm"    "gyros_arm_x"
## [19] "gyros_arm_y"      "gyros_arm_z"        "accel_arm_x"
## [22] "accel_arm_y"      "accel_arm_z"        "magnet_arm_x"
## [25] "magnet_arm_y"     "magnet_arm_z"       "roll_dumbbell"
## [28] "pitch_dumbbell"   "yaw_dumbbell"       "total_accel_dumbbell"
## [31] "gyros_dumbbell_x" "gyros_dumbbell_y"   "gyros_dumbbell_z"
## [34] "accel_dumbbell_x" "accel_dumbbell_y"   "accel_dumbbell_z"
## [37] "magnet_dumbbell_x" "magnet_dumbbell_y"  "magnet_dumbbell_z"
## [40] "roll_forearm"     "pitch_forearm"      "yaw_forearm"
## [43] "total_accel_forearm" "gyros_forearm_x"    "gyros_forearm_y"
## [46] "gyros_forearm_z"   "accel_forearm_x"     "accel_forearm_y"
## [49] "accel_forearm_z"   "magnet_forearm_x"    "magnet_forearm_y"
## [52] "magnet_forearm_z"  "classe"
```

We also check if any of the variables have near zero variance.

```
mydata <- nearZeroVar(trained, saveMetrics = TRUE)
mydata
```

```
##               freqRatio percentUnique zeroVar  nzv
## roll_belt      1.101904      6.7781062  FALSE FALSE
## pitch_belt     1.036082      9.3772296  FALSE FALSE
## yaw_belt       1.058480      9.9734991  FALSE FALSE
## total_accel_belt 1.063160      0.1477933  FALSE FALSE
## gyros_belt_x    1.058651      0.7134849  FALSE FALSE
## gyros_belt_y    1.144000      0.3516461  FALSE FALSE
## gyros_belt_z    1.066214      0.8612782  FALSE FALSE
## accel_belt_x    1.055412      0.8357966  FALSE FALSE
## accel_belt_y    1.113725      0.7287738  FALSE FALSE
## accel_belt_z    1.078767      1.5237998  FALSE FALSE
## magnet_belt_x   1.090141      1.6664968  FALSE FALSE
## magnet_belt_y   1.099688      1.5187035  FALSE FALSE
## magnet_belt_z   1.006369      2.3290184  FALSE FALSE
## roll_arm       52.338462     13.5256345  FALSE FALSE
## pitch_arm      87.256410     15.7323412  FALSE FALSE
## yaw_arm        33.029126     14.6570176  FALSE FALSE
## total_accel_arm 1.024526      0.3363572  FALSE FALSE
## gyros_arm_x     1.015504      3.2769341  FALSE FALSE
## gyros_arm_y     1.454369      1.9162165  FALSE FALSE
## gyros_arm_z     1.110687      1.2638875  FALSE FALSE
## accel_arm_x     1.017341      3.9598410  FALSE FALSE
## accel_arm_y     1.140187      2.7367241  FALSE FALSE
## accel_arm_z     1.128000      4.0362858  FALSE FALSE
## magnet_arm_x    1.000000      6.8239731  FALSE FALSE
## magnet_arm_y    1.056818      4.4439914  FALSE FALSE
## magnet_arm_z    1.036364      6.4468454  FALSE FALSE
## roll_dumbbell   1.022388     84.2065029  FALSE FALSE
## pitch_dumbbell  2.277372     81.7449801  FALSE FALSE
## yaw_dumbbell    1.132231     83.4828254  FALSE FALSE
## total_accel_dumbbell 1.072634      0.2191418  FALSE FALSE
## gyros_dumbbell_x 1.003268      1.2282132  FALSE FALSE
## gyros_dumbbell_y 1.264957      1.4167771  FALSE FALSE
## gyros_dumbbell_z 1.060100      1.0498420  FALSE FALSE
```

```
## accel_dumbbell_x      1.018018      2.1659362    FALSE FALSE
## accel_dumbbell_y      1.053061      2.3748853    FALSE FALSE
## accel_dumbbell_z      1.133333      2.0894914    FALSE FALSE
## magnet_dumbbell_x     1.098266      5.7486495    FALSE FALSE
## magnet_dumbbell_y     1.197740      4.3012945    FALSE FALSE
## magnet_dumbbell_z     1.020833      3.4451126    FALSE FALSE
## roll_forearm          11.589286     11.0895933    FALSE FALSE
## pitch_forearm         65.983051     14.8557741    FALSE FALSE
## yaw_forearm           15.322835     10.1467740    FALSE FALSE
## total_accel_forearm    1.128928      0.3567424    FALSE FALSE
## gyros_forearm_x        1.059273      1.5187035    FALSE FALSE
## gyros_forearm_y        1.036554      3.7763735    FALSE FALSE
## gyros_forearm_z        1.122917      1.5645704    FALSE FALSE
## accel_forearm_x        1.126437      4.0464784    FALSE FALSE
## accel_forearm_y        1.059406      5.1116094    FALSE FALSE
## accel_forearm_z        1.006250      2.9558659    FALSE FALSE
## magnet_forearm_x       1.012346      7.7667924    FALSE FALSE
## magnet_forearm_y       1.246914      9.5403119    FALSE FALSE
## magnet_forearm_z       1.000000      8.5771073    FALSE FALSE
## classe                 1.469581      0.0254816    FALSE FALSE
```

Since all the variables return FALSE, we are left with 52 variables to build our model. Before this we apply the same cleaning procedures to the final testing dataset with 20 cases.

```
test_cleaned <- testing[, colSums(is.na(testing)) == 0]
remove_test_col <- grepl("~X|user_name|timestamp|window",names(test_cleaned))
test_cleaned <- test_cleaned[,!remove_test_col]
testClean <- test_cleaned[, sapply(test_cleaned, is.numeric)]
dim(testClean)
```

```
## [1] 20 53
```

```
dim(trClean)
```

```
## [1] 19622    52
```

We should also make sure the training and test datasets have same columns/names.

```
names(test_cleaned) == names(trained)
```

```
## [1] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
## [12] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
## [23] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
## [34] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
## [45] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE FALSE
```

Last column names is different, but that is OK.

Building predictions

We split data into training and testing samples, and set a unique seed so we can reproduce the results. We will use 65 35 split for training and testing.

```
set.seed(83912)
inTrain <- createDataPartition(y=trained$clas,p=0.65, list=FALSE)
my_training <- trained[inTrain,]
my_testing <- trained[-inTrain,]
```

Predicting with decision trees

```
model_Dtree <- rpart(classe ~.,data = my_training, method = "class")
prediction_Dtree <- predict(model_Dtree, my_testing, type = "class")
confusionMatrix(prediction_Dtree, my_testing$classe)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    A    B    C    D    E
##           A 1776  210   22   49   20
##           B   55  721   48   91   83
##           C   44  197  978  177  152
##           D   21  101   68  705   62
##           E   57   99   81  103  945
##
## Overall Statistics
##
##           Accuracy : 0.7465
##           95% CI : (0.7361, 0.7568)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.6789
##           McNemar's Test P-Value : < 2.2e-16
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9094  0.5429  0.8170  0.6267  0.7488
## Specificity      0.9387  0.9500  0.8994  0.9561  0.9393
## Pos Pred Value   0.8551  0.7224  0.6318  0.7367  0.7354
## Neg Pred Value   0.9630  0.8965  0.9588  0.9289  0.9432
## Prevalence       0.2845  0.1934  0.1744  0.1639  0.1838
## Detection Rate   0.2587  0.1050  0.1425  0.1027  0.1377
## Detection Prevalence 0.3025  0.1454  0.2255  0.1394  0.1872
## Balanced Accuracy 0.9240  0.7464  0.8582  0.7914  0.8441
```

```
model_Dtree
```

```
## n= 12757
```

```

##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
##      1) root 12757 9130 A (0.28 0.19 0.17 0.16 0.18)
##          2) roll_belt< 129.5 11642 8057 A (0.31 0.21 0.19 0.18 0.11)
##              4) pitch_forearm< -33.95 1005      6 A (0.99 0.006 0 0 0) *
##              5) pitch_forearm>=-33.95 10637 8051 A (0.24 0.23 0.21 0.2 0.12)
##                  10) magnet_dumbbell_y< 438.5 8968 6442 A (0.28 0.18 0.24 0.19 0.11)
##                      20) roll_forearm< 123.5 5610 3360 A (0.4 0.19 0.18 0.17 0.061)
##                          40) magnet_dumbbell_z< -27.5 1874  620 A (0.67 0.21 0.012 0.078 0.034)
##                              80) roll_forearm>=-136.5 1554  333 A (0.79 0.17 0.013 0.028 0.0058) *
##                              81) roll_forearm< -136.5 320  193 B (0.1 0.4 0.0094 0.32 0.17) *
##                                  41) magnet_dumbbell_z>=-27.5 3736 2740 A (0.27 0.17 0.26 0.22 0.074)
##                                      82) accel_dumbbell_y>=-40.5 3327 2334 A (0.3 0.19 0.18 0.24 0.083)
##                                          164) yaw_belt>=168.5 484   74 A (0.85 0.087 0.0021 0.06 0.0041) *
##                                          165) yaw_belt< 168.5 2843 2058 D (0.21 0.21 0.21 0.28 0.096)
##                                              330) pitch_belt< -43.15 302   36 B (0.0066 0.88 0.07 0.02 0.023) *
##                                              331) pitch_belt>=-43.15 2541 1762 D (0.23 0.13 0.23 0.31 0.1)
##                                                  662) roll_belt>=125.5 546  226 C (0.35 0.049 0.59 0.011 0)
##                                                      1324) magnet_belt_z< -322.5 163    3 A (0.98 0.0061 0.012 0 0) *
##                                                      1325) magnet_belt_z>=-322.5 383   65 C (0.086 0.068 0.83 0.016 0) *
##                                                          663) roll_belt< 125.5 1995 1222 D (0.19 0.15 0.13 0.39 0.13)
##                                                              1326) yaw_belt< -85.55 1241  948 A (0.24 0.21 0.14 0.22 0.19)
##                                                                  2652) accel_dumbbell_z< 25.5 769  490 A (0.36 0.14 0.22 0.24 0.033)
##                                                                      5304) yaw_forearm>=-94.55 569  290 A (0.49 0.18 0.24 0.053 0.032)
##                                                                          10608) magnet_forearm_z>=-154.5 358   84 A (0.77 0.14 0.02 0.039 0.034) *
##                                                                          10609) magnet_forearm_z< -154.5 211   80 C (0.024 0.25 0.62 0.076 0.028) *
##                                                                              5305) yaw_forearm< -94.55 200   44 D (0 0.025 0.16 0.78 0.035) *
##                                                                                  2653) accel_dumbbell_z>=25.5 472  264 E (0.03 0.33 0.013 0.19 0.44)
##                                                                                      5306) roll_dumbbell< 38.68377 157   32 B (0.032 0.8 0.038 0.025 0.11) *
##                                                                                      5307) roll_dumbbell>=38.68377 315  124 E (0.029 0.098 0 0.27 0.61) *
##                                                                                          1327) yaw_belt>=-85.55 754  255 D (0.13 0.048 0.12 0.66 0.044)
##                                                                                              2654) yaw_arm< -112.35 92    0 A (1 0 0 0 0) *
##                                                                                                  2655) yaw_arm>=-112.35 662  163 D (0.0045 0.054 0.14 0.75 0.05) *
##                                                                                                      83) accel_dumbbell_y< -40.5 409   36 C (0.0073 0.042 0.91 0.034 0.0049) *
##                                                                                                          21) roll_forearm>=123.5 3358 2216 C (0.082 0.17 0.34 0.22 0.18)
##                                                                                                              42) magnet_dumbbell_y< 291.5 1988 1002 C (0.1 0.12 0.5 0.15 0.14)
##                                                                                                                  84) magnet_forearm_z< -251.5 157   31 A (0.8 0.07 0 0.038 0.089) *
##                                                                                                                  85) magnet_forearm_z>=-251.5 1831  845 C (0.039 0.13 0.54 0.16 0.14) *
##                                                                                                  43) magnet_dumbbell_y>=291.5 1370  913 D (0.057 0.24 0.11 0.33 0.26)
##                                                                                                      86) accel_forearm_x>=-100.5 880  572 E (0.051 0.31 0.15 0.14 0.35)
##                                                                                                          172) roll_dumbbell< 41.02769 180   36 B (0.056 0.8 0.011 0.039 0.094) *
##                                                                                                          173) roll_dumbbell>=41.02769 700  409 E (0.05 0.18 0.19 0.16 0.42) *
##                                                                                          87) accel_forearm_x< -100.5 490  153 D (0.067 0.12 0.041 0.69 0.088) *
##                                                                                              11) magnet_dumbbell_y>=438.5 1669  817 B (0.036 0.51 0.047 0.22 0.19)
##                                                                                                  22) total_accel_dumbbell>=5.5 1218  439 B (0.049 0.64 0.063 0.021 0.23)
##                                                                                                      44) roll_belt>=-0.58 1030  251 B (0.058 0.76 0.075 0.024 0.086) *
##                                                                                                      45) roll_belt< -0.58 188    0 E (0 0 0 0 1) *
##                                                                                                          23) total_accel_dumbbell< 5.5 451  108 D (0 0.16 0.0022 0.76 0.075) *
##                                                                                          3) roll_belt>=129.5 1115   42 E (0.038 0 0 0 0.96) *

```

I find that decision trees are not good enough to get the sort of accuracy we are after. Therefore I will try random forests now. ### Predicting with random forests

```

model_RF1 <- randomForest(classe ~., data = my_training)
prediction_RF1 <- predict(model_RF1, my_testing, type = "class")
confusionMatrix(prediction_RF1, my_testing$classe)

```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction      A      B      C      D      E
```

```
##           A 1951      3      0      1      0
```

```
##           B   2 1320     13      0      0
```

```
##           C    0      5 1184     15      2
```

```
##           D    0      0      0 1108      1
```

```
##           E    0      0      0      1 1259
```

```
##
```

```
## Overall Statistics
```

```
##
```

```
##           Accuracy : 0.9937
```

```
##           95% CI : (0.9916, 0.9955)
```

```
##           No Information Rate : 0.2845
```

```
##           P-Value [Acc > NIR] : < 2.2e-16
```

```
##
```

```
##           Kappa : 0.9921
```

```
##           McNemar's Test P-Value : NA
```

```
##
```

```
## Statistics by Class:
```

```
##
```

```
##           Class: A Class: B Class: C Class: D Class: E
```

```
## Sensitivity          0.9990    0.9940    0.9891    0.9849    0.9976
```

```
## Specificity          0.9992    0.9973    0.9961    0.9998    0.9998
```

```
## Pos Pred Value       0.9980    0.9888    0.9818    0.9991    0.9992
```

```
## Neg Pred Value       0.9996    0.9986    0.9977    0.9970    0.9995
```

```
## Prevalence           0.2845    0.1934    0.1744    0.1639    0.1838
```

```
## Detection Rate       0.2842    0.1923    0.1725    0.1614    0.1834
```

```
## Detection Prevalence 0.2848    0.1945    0.1757    0.1615    0.1835
```

```
## Balanced Accuracy     0.9991    0.9956    0.9926    0.9924    0.9987
```

```
model_RF1
```

```
##
```

```
## Call:
```

```
## randomForest(formula = classe ~ ., data = my_training)
```

```
##           Type of random forest: classification
```

```
##           Number of trees: 500
```

```
##           No. of variables tried at each split: 7
```

```
##
```

```
##           OOB estimate of error rate: 0.55%
```

```
## Confusion matrix:
```

```
##           A      B      C      D      E class.error
```

```
## A 3621      5      1      0      0 0.001654260
```

```
## B  12 2453      4      0      0 0.006480356
```

```
## C   0  13 2209      3      0 0.007191011
```

```
## D   0   0  20 2069      2 0.010521282
```

```
## E   0   0   3   7 2335 0.004264392
```

When running RF with 500 trees we get a very good accuracy of about 99.4 percent. I find that this does not vary much if we reduce the number of trees to 100 (Acc = 0.9932). Adding k-fold cross-validation does not improve the result. Pre-processing (centering and scaling) also does not change the result. ### Adding cross validation

```
control <- trainControl(method = "cv",5)
model_RF2 <- train(classe ~., data = my_training, method = "rf",trControl=control, ntree=100)
prediction_RF2 <- predict(model_RF2, my_testing)
confusionMatrix(prediction_RF2, my_testing$classe)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction      A      B      C      D      E
##           A 1949      6      0      0      0
##           B   4 1315     10      1      2
##           C    0      7 1181     13      5
##           D    0      0      6 1111      0
##           E    0      0      0      0 1255
```

```
##
```

```
## Overall Statistics
```

```
##
```

```
##           Accuracy : 0.9921
```

```
##           95% CI : (0.9897, 0.9941)
```

```
##           No Information Rate : 0.2845
```

```
##           P-Value [Acc > NIR] : < 2.2e-16
```

```
##
```

```
##           Kappa : 0.99
```

```
##           McNemar's Test P-Value : NA
```

```
##
```

```
## Statistics by Class:
```

```
##
```

```
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9980  0.9902  0.9866  0.9876  0.9945
## Specificity      0.9988  0.9969  0.9956  0.9990  1.0000
## Pos Pred Value    0.9969  0.9872  0.9793  0.9946  1.0000
## Neg Pred Value    0.9992  0.9977  0.9972  0.9976  0.9988
## Prevalence        0.2845  0.1934  0.1744  0.1639  0.1838
## Detection Rate    0.2839  0.1916  0.1720  0.1618  0.1828
## Detection Prevalence 0.2848  0.1940  0.1757  0.1627  0.1828
## Balanced Accuracy 0.9984  0.9936  0.9911  0.9933  0.9972
```

```
model_RF2
```

```
## Random Forest
```

```
##
```

```
## 12757 samples
```

```
##    52 predictors
```

```
##    5 classes: 'A', 'B', 'C', 'D', 'E'
```

```
##
```

```
## No pre-processing
```

```
## Resampling: Cross-Validated (5 fold)
```

```
## Summary of sample sizes: 10206, 10205, 10205, 10206, 10206
```



```
## Resampling results across tuning parameters:
##
##   mtry  Accuracy   Kappa
##    2    0.9891044 0.9862154
##   27    0.9909853 0.9885956
##   52    0.9830681 0.9785807
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 27.
```

I can work out the accuracy and out of sample error on the RF1 model, which I also apply to the final test set.

```
accuracy <- postResample(prediction_RF1, my_testing$classe)
accuracy
```

```
## Accuracy      Kappa
## 0.9937363 0.9920766
```

```
OutSampleErr <- 1 - as.numeric(confusionMatrix(prediction_RF1, my_testing$classe)$overall[1])
OutSampleErr
```

```
## [1] 0.006263656
```

Running this on the test set for the quiz


```
fin_results <- predict(model_RF1, test_cleaned[, -length(names(test_cleaned))])
fin_results
```

```
##  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20
##  B  A  B  A  A  E  D  B  A  A  B  C  B  A  E  E  A  B  B  B
## Levels: A B C D E
```

Plots


```
plot(model_RF1)
```

model_RF1

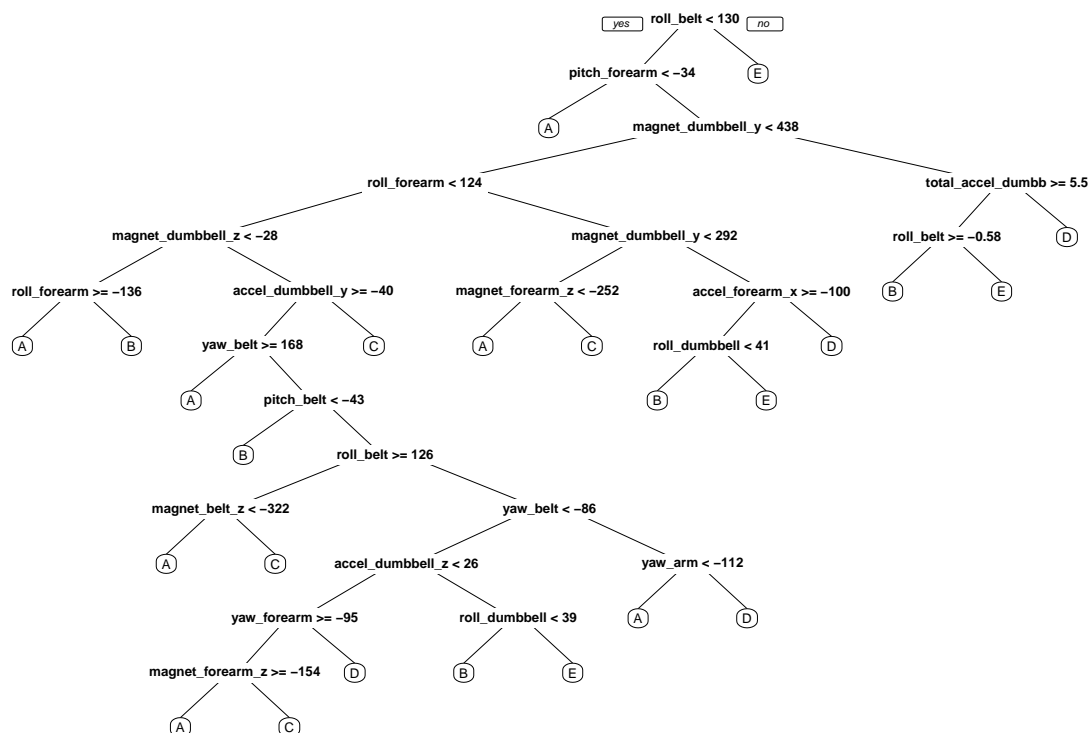


Correlation matrix

```
corrPlot <- cor(my_training[, -length(names(my_training))])  
corrplot(corrPlot, method="color")
```




```
prp(model_Dtree)
```



```
fancyRpartPlot(model_Dtree)
```

```
## Warning: labs do not fit even at cex 0.15, there may be some overplotting
```



Rattle 2016-Jun-15 21:27:47 zablocki